

Introduction to Artificial Intelligence



VICTORIA UNIVERSITY OF
WELLINGTON
TE HERENGA WAKA

COMP307/AIML420

Neural Networks: Tutorial

Dr Fangfang Zhang

fangfang.zhang@ecs.vuw.ac.nz

COMP307/AIML420 Week 4 (Tutorial)


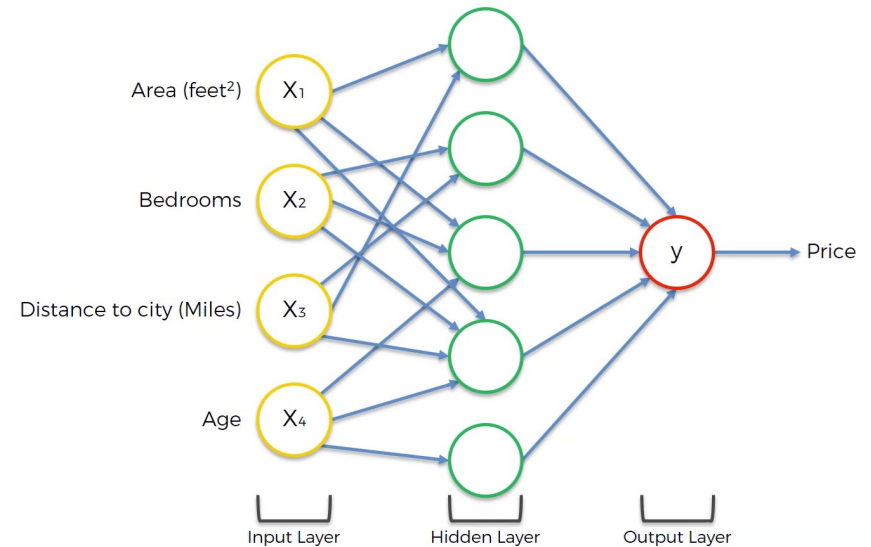

1. Announcements

- Assignment 1 (**15%**)
- Due on 30th Mar.
- Helpdesk (2 hours)
this Thurs.---next Wed.
- Part 1 and Part 2
- Part 3

2. Neural Networks

- Perceptron (Part 3)
- Back Propagation

Neural Networks



- Flexible structure
- Number of inputs
- Number of layers
- Fully/partially connected
- Number of outputs

Depends on the problems!

Ups and downs of Neural Networks

- 1958: Perceptron
- 1969: Perceptron has limitations
- 1980s: Multi-layer perceptron

Do not have a significant difference from DNN today

- 1986: Backpropagation

Improve the efficiency of NN learning

- 1989: 1 hidden layer is “good enough”, why deep?
- 2011: start to be popular
-

Perceptron

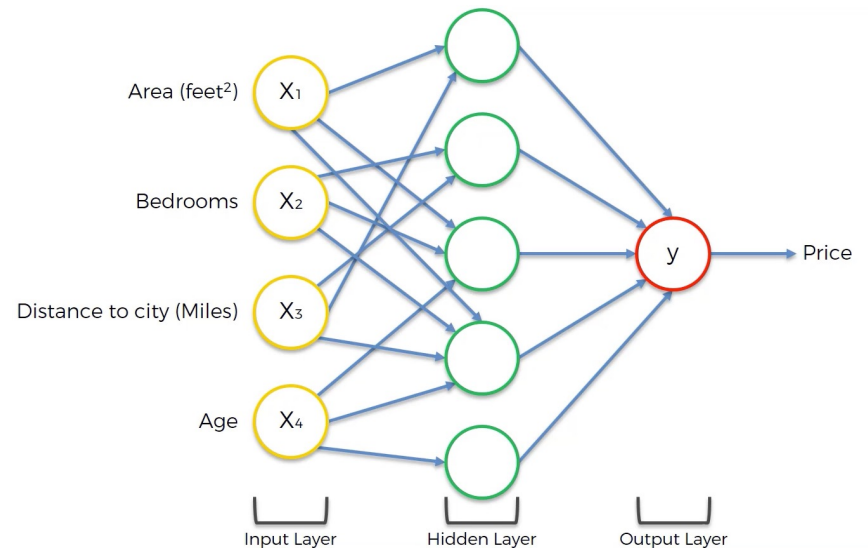
- Three steps for learning



- neural networks
- decision tree
- accuracy

- Different dataset

- image data
- text data
- speech data
- tabular data (with numbers)

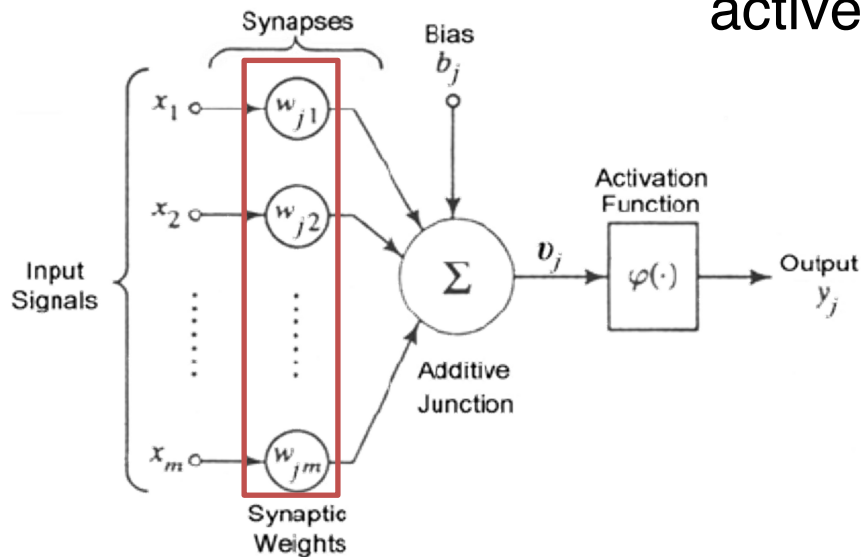


Perceptron


- What we have (classification task):

Dataset (features, labels)

Neural network structure (features=input, labels=output, active function)



$$y_j = \begin{cases} 1, & \text{if } \sum_{i=1}^m w_{ji}x_i + b_j > 0, \\ 0, & \text{otherwise} \end{cases}$$




Active function

- Perceptron learning
--- learning weights and bias

Perceptron

- How to get the **optimal weights** and **bias**?
--- important features may have larger weights
- To simplify notation, we can transform the bias to a **weight** $w_{j0} = b_j$, where $x_0 = 1$ (**dummy feature**) always holds

$$y_j = \begin{cases} 1, & \text{if } \sum_{i=1}^m w_{ji}x_i + b_j > 0, \\ 0, & \text{otherwise} \end{cases}$$



$$b_j = w_{j0} \cdot 1 = w_{j0}x_0$$

$$l_j = \begin{cases} 1, & \text{if } \sum_{i=0}^m w_{ji}x_i > 0, \\ 0, & \text{otherwise} \end{cases}$$

Perceptron

1. Define a dummy feature (for bias)
2. **Random initialise** a set of weights (how many?)
3. Get the first instance in the dataset
4. **Sum up** *feature values and weights*, and get predicted class label along with active function

$$y_j = \begin{cases} 1, & \text{if } \sum_{i=0}^m w_{ji}x_i > 0, \\ 0, & \text{otherwise} \end{cases}$$

5. **Adjust** the weights (class labels, i.e., 1, 0)

$$w_i \leftarrow w_i + (d - y)x_i, i = 0, 1, 2, \dots, m$$


d is the desired class label, y is the predicted class label

- $d = 1, y = 1$ or $d = 0, y = 0$ (same) *nothing*
- $d = 1, y = 0, d - y > 0$, **increase** (different)
- $d = 0, y = 1, d - y < 0$, **decrease** (different)

6. Go to step 4, and use the next instance (until meet the stop criterion)

dataset


length	width	weight	class
12	7	43	A (1)
1	24	51	B (0)



NN Example

- Calculate the outputs of this network (feedforward): to 2dp

I_1	I_2	w_{13}	w_{14}	w_{23}	w_{24}	w_{35}	w_{36}	w_{45}	w_{46}	b_3	b_4	b_5	b_6
0.90	-0.20	0.72	-0.31	0.10	-0.92	-0.37	0.43	-0.19	0.78	0.01	0.38	-0.13	0.78



$$\begin{aligned}
 Z_3 &= w_{13} * I_1 + w_{23} * I_2 + b_3 \\
 &= 0.72 * 0.90 + 0.10 * (-0.2) + 0.01 \\
 &= 0.64
 \end{aligned}$$

$$\begin{aligned}
 Z_4 &= w_{14} * I_1 + w_{24} * I_2 + b_4 \\
 &= (-0.31) * 0.90 + (-0.92) * (-0.2) \\
 &\quad + 0.38 \\
 &= 0.29
 \end{aligned}$$

$$z_j = \sum_i w_{ji} x_i + b_j$$


z_3	0.64
O_3	0.65
z_4	0.29
O_4	0.57
z_5	-0.48
O_5	0.62
z_6	1.50
O_6	0.82

- Weighted sum of a node:
- Output of a node:
 - Where φ is the activation function

3. Assume φ is the sigmoid function: $O_j = \varphi(z_j) = \frac{1}{1+e^{-z_j}}$

Class = ? B

NNs for (Multi-Class) Classification



Training a Neural Network


- **Initialise** the weights (randomly)
- **Feedforward**
 - For each example/instance, calculate the **predicted outputs** o_z using the current weights
 - Calculate the total **error** $\sum_z (d_z - o_z)^2$
 - d_z means “*desired*”
 - o_z means “output” (i.e. what we actually got)
- If the error is small enough, we can stop.
- Otherwise, we use **back propagation** to adjust the weights to make the error *smaller*.
 - Uses gradient descent (GD)

Back Propagation (BP) Algorithm


- **Improve the efficiency of NN learning**
- **How to update the weights**
- Estimate the contribution (gradient) of each weight to the *error*, i.e. how much the error will be reduced by changing the weight (gradient)
- **Change** each weight (**simultaneously**) proportional to its **contribution** to reduce the error as **much as possible**
 - Move in the direction of the **steepest gradient**
- We calculate the contribution/gradient **backwards** (from the last/output layer to the first hidden layer)

Notes on BP Algorithm

- *1 Epoch*: all input examples (entire training set, batch, ...)
- A target of 0 or 1 cannot reasonably be reached. Usually interpret an output > 0.9 or > 0.8 as '1'
- Training may require *thousands* of epochs. A convergence curve will help to decide when to stop (over-fitting?)




Back Propagation




This is a function.
 Input vector, output vector $f\left(\begin{bmatrix} 1 \\ -1 \end{bmatrix}\right) = \begin{bmatrix} 0.62 \\ 0.83 \end{bmatrix}$ $f\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}\right) = \begin{bmatrix} 0.51 \\ 0.85 \end{bmatrix}$

Back Propagation

Input




Output



Back Propagation

Total Loss

For all training data ...



Total Loss:

$$L = \sum_{n=1}^N C^n$$

Find a function in function set that minimizes total loss L

Find the network parameters θ^* that minimize total loss L

Back Propagation


- Chain Rule

Case 1 $y = g(x) \quad z = h(y)$

$$\Delta x \rightarrow \Delta y \rightarrow \Delta z \qquad \frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$$

Case 2



$$x = g(s) \qquad y = h(s) \qquad z = k(x, y)$$



$$\frac{dz}{ds} = \frac{\partial z}{\partial x} \frac{dx}{ds} + \frac{\partial z}{\partial y} \frac{dy}{ds}$$

Back Propagation (BP) Algorithm

- How **big a change** should we make to **weight $w_{i \rightarrow j}$** ?
 - Make a **big change** if will improve error **a lot** (big contribution)
 - Make a **small change** if there is **little effect** on error (small contribution)



- β_j is how “**beneficial**” a change is for node j (“error term”)
- When changing $w_{i \rightarrow j}$, the error change should be:
 - Proportional to the **output**: o_i (larger output = more effect)
 - Proportional to the **slope of the activation function** at node j : slope_j
 - Proportional to error term of j (β_j)

BP Algorithm Implementation

- Initialise all weights (+bias) to **small random values**
- Until total error is small enough, repeat:
 - For each input example:
 - **Feed forward pass** to get predicted outputs
 - Compute $\beta_z = d_z - o_z$ for each output node
 - Compute $\beta_j = \sum_k w_{j \rightarrow k} o_k (1 - o_k) \beta_k$ for each hidden node (working backwards from last to first layer)
 - Compute (+store) the weight changes for all weights
$$\Delta w_{i \rightarrow j} = \eta o_i o_j (1 - o_j) \beta_j$$
(proportional to all 3 factors), Let η be the learning rate
 - Sum up weight changes for all input examples
 - Change weights!

NN Example: Your Turn!


- Calculate the new weights and biases (backprop): to 2dp

d_5	d_6	η	β_3	β_4	β_5	β_6
0	1	1				

Inputs

Hidden

Outputs



w_{13}	
w_{14}	
w_{23}	
w_{24}	
w_{35}	
w_{36}	
w_{45}	
w_{46}	
b_3	
b_4	
b_5	
b_6	

Useful Formulae: Backprop

- Error term of an output node: $\beta_j = d_j - O_j$
- Error term of a hidden node: $\beta_j = \sum_k w_{j \rightarrow k} O_k (1 - O_k) \beta_k$
 - (For the sigmoid activation function)
- Amount to change a weight: $\Delta w_{i \rightarrow j} = \eta O_i O_j (1 - O_j) \beta_j$
- Amount to change a bias: $\Delta b_j = \eta O_j (1 - O_j) \beta_j$

Summary

- Perceptron
- Back Propagation
- Next week
 - Neural Engineering (next Monday)
 - Evolutionary Computation (next Tuesday)