

Introduction to Artificial Intelligence



VICTORIA UNIVERSITY OF
WELLINGTON
TE HERENGA WAKA

COMP307/AIML420

Neural Networks 1: The Perceptron

Dr Andrew Lensen



Andrew.Lensen@vuw.ac.nz

Outline

- Why ANN?
- Origin
- The perceptron
- Perceptron learning
- What can (not) perceptron learn
- Extending perceptron

Why Artificial Neural Networks (ANNs)?





- “Killer” applications in a lot of areas
 - Computer vision/image processing
 - Playing games (AlphaGo, Watson, ...)
 - Big data
 - ...



Origin


- Human brain shows amazing capability in:
 - Learning
 - Perception
 - Adaptability
 - ...
- Can we *simulate the* human brain to achieve the above functionalities?
- **Artificial** neural networks use this as inspiration/motivation

Origin

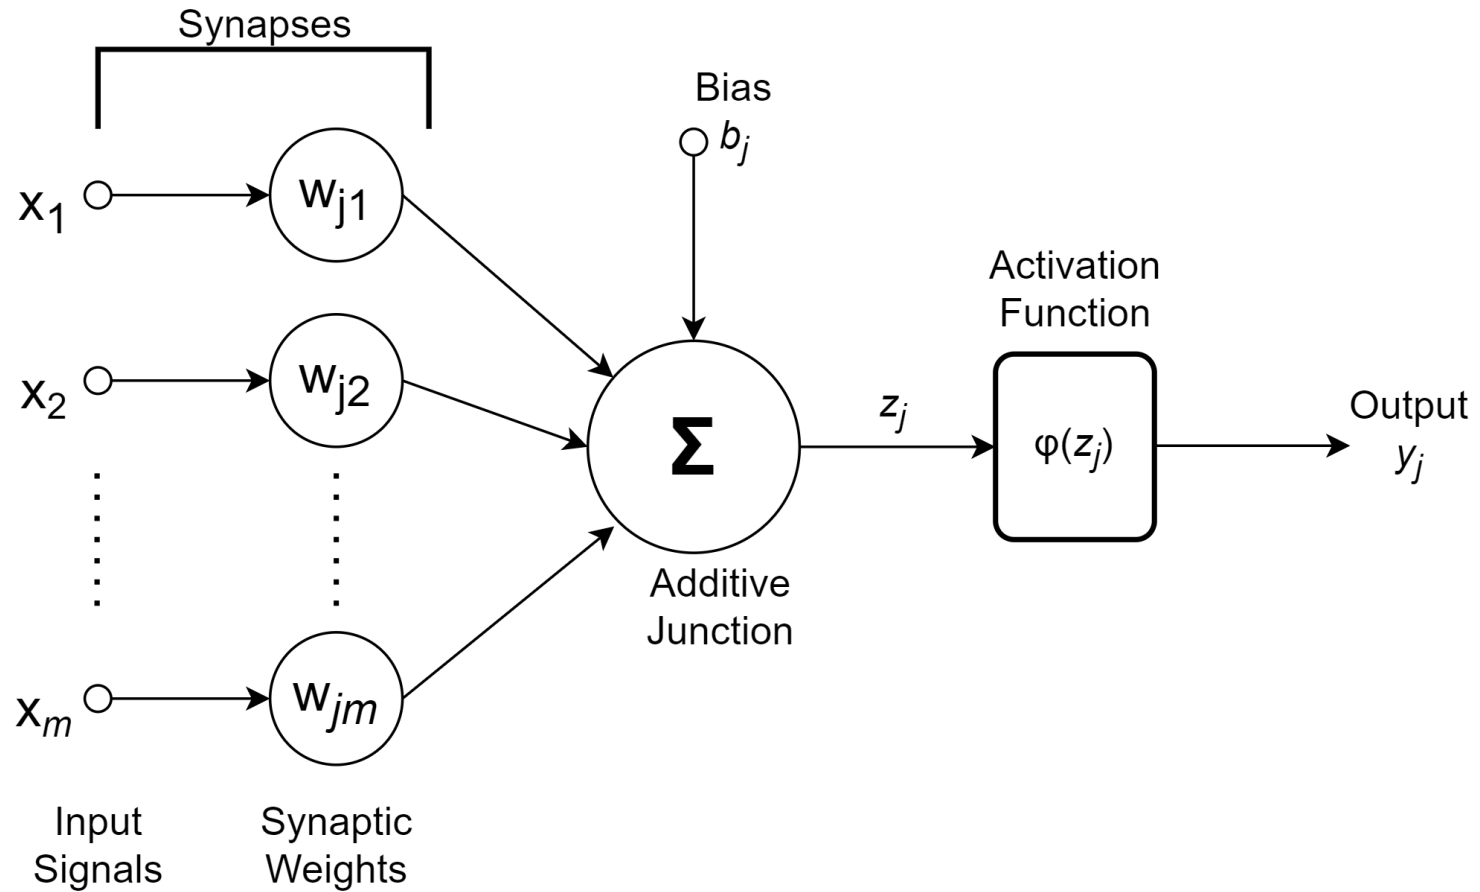


Origin

- Facts about human brain:
 - 10^{11} neurons, massively connected
 - Each neuron is connected to 1000–10,000 other neurons
 - 10^{14} to 10^{15} connections in total!
 - Brain message passing is 1 million times **slower** than modern electronic circuits
 - But very *efficient* for complex decision making
 - Usually fewer than 100 serial stages
 - 100 step rule (500ms response time)




A Single Artificial Neuron



$$z_j = \sum_{i=1}^m w_{ji} x_i + b_j$$


$$y_j = \varphi(z_j)$$

Activation Functions



Threshold

$$y_j = \begin{cases} 1, & \text{if } z_j > 0, \\ 0, & \text{otherwise} \end{cases}$$




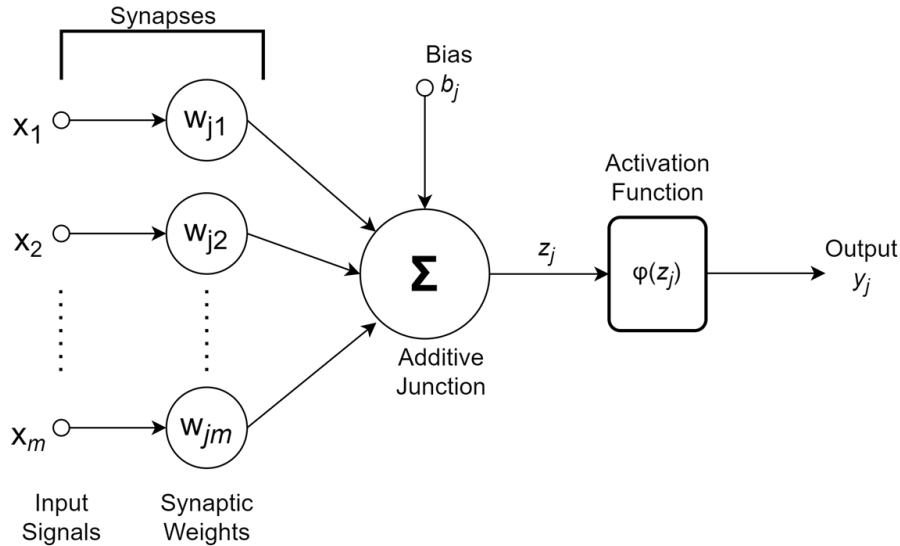
Sigmoid

$$y_j = \frac{1}{1 + e^{-z_j}}$$

Why do we need this?

Perceptron

- A *special* type of artificial neuron
 - Real-valued inputs
 - Binary output
 - Threshold activation function







$$y = \begin{cases} 1, & \text{if } x_1 + x_2 > 0, \\ 0, & \text{otherwise.} \end{cases}$$

$$y_j = \begin{cases} 1, & \text{if } \sum_{i=1}^m w_{ji}x_i + b_j > 0, \\ 0, & \text{otherwise} \end{cases}$$

Perceptron

- To perform **linear** classification
 - 2 inputs: a **line**; 3 inputs: a **plane**
- Can do **online** learning
 - Update w_{ji} and b_j based on new examples



Learning A Perceptron

- How to get the **optimal weights** and **bias**?
- Let us only consider **accuracy**:
 - *Optimal* if **100% accuracy** on training set
 - Can have **many optimal** solutions...
- To simplify notation, we can transform the bias to a **weight** $w_{j0} = b_j$, where $x_0 = 1$ always holds

$$y_j = \begin{cases} 1, & \text{if } \sum_{i=1}^m w_{ji}x_i + b_j > 0, \\ 0, & \text{otherwise} \end{cases}$$

$$b_j = w_{j0} \cdot 1 = w_{j0}x_0$$

$$y_j = \begin{cases} 1, & \text{if } \sum_{i=0}^m w_{ji}x_i > 0, \\ 0, & \text{otherwise} \end{cases}$$

Learning A Perceptron

- Initialise weights and threshold randomly (or set all to zero)
 - Given a new example/instance $(x_1, x_2, \dots, x_m, d)$:
 - Input feature vector: (x_1, x_2, \dots, x_m)
 - Output (desired class label): d
 - Predicted output: y
- $$y = \begin{cases} 1, & \text{if } \sum_{i=0}^m w_i x_i > 0, \\ 0, & \text{otherwise} \end{cases}$$
- If $y = 0$ and $d = 1$,
 - increase $b = w_0$, increase w_i for each positive x_i ,
decrease w_i for each negative x_i
 - If $y = 1$ and $d = 0$,
 - decrease $b = w_0$, decrease w_i for each positive x_i ,
increase w_i for each negative x_i
 - Repeat for each new example until desired result achieved

Learning A Perceptron

- Initialise weights and threshold randomly (or set all to zero)
- Given a new example/instance $(x_1, x_2, \dots, x_m, d)$:
 - **Input** feature vector: (x_1, x_2, \dots, x_m)
 - **Output** (desired class label): d
 - **Predicted** output: y

$$y = \begin{cases} 1, & \text{if } \sum_{i=0}^m w_i x_i > 0, \\ 0, & \text{otherwise} \end{cases}$$


- We can use the equation:

$$w_i \leftarrow w_i + \eta(d - y)x_i, i = 0, 1, 2, \dots, m$$

- Where $\eta \in [0,1]$ is called the **learning rate**
- **Repeat** for each new example until desired result achieved

Problem with the Perceptron

- What can the perceptron learn?





- Perceptron convergence theorem:* The perceptron learning algorithm will converge **if and only if** the training set is **linearly separable**.
- Cannot learn for XOR (Minsky and Papert, 1969)

Multi-Layer Perceptron (MLP)


- Add **one hidden node** between the inputs and output

x1	x2	y (class)
0	0	0
1	0	1
0	1	1
1	1	0



Neural Network

- Add **more hidden layers**
- Add **more nodes** for hidden layers



Summary

- ANN has many “killer” applications
 - Image analysis, playing games, ...
- Perceptron – the simplest neural network
- How to learn a perceptron
- Limitations of perceptron, multi-layer perceptron, general neural network
- Reading: Textbook section 20.5 (2nd edition) or section 18.7 (3rd/4th edition)