

Introduction to Artificial Intelligence



VICTORIA UNIVERSITY OF
WELLINGTON
TE HERENGA WAKA

COMP307/AIML420

Neural Networks 2: Back Propagation

Dr Andrew Lensen

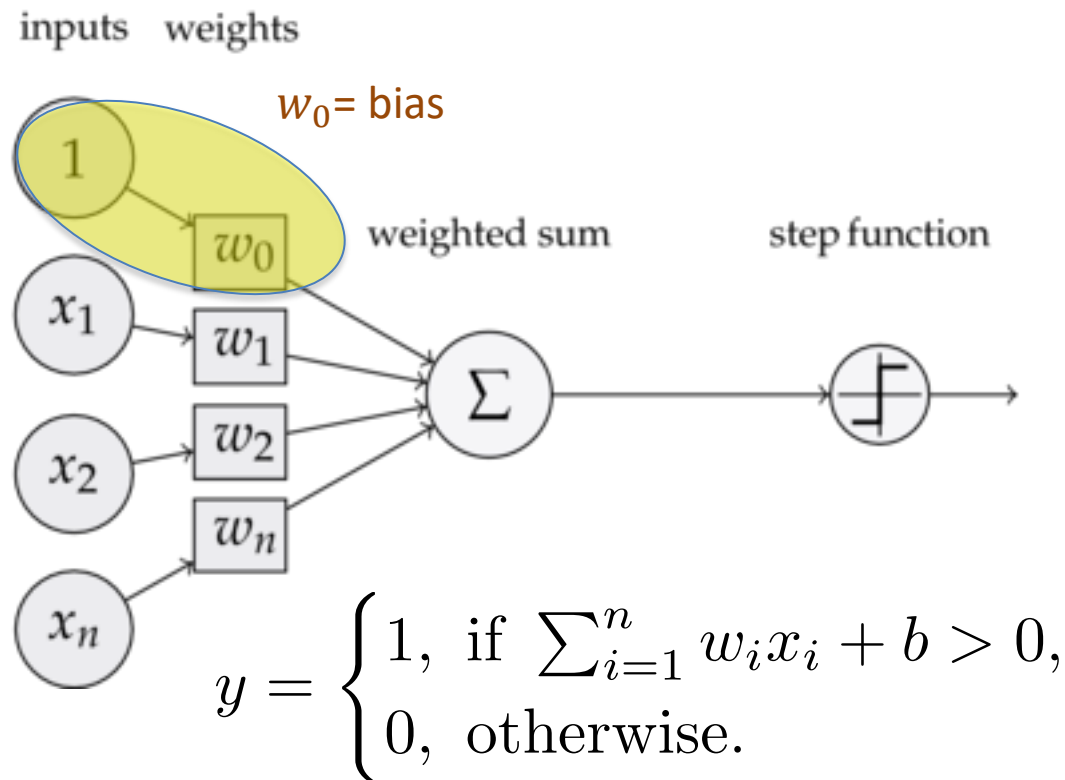
Andrew.Lensen@vuw.ac.nz

Outline

- Revisiting (Multi-layer) Perceptron
- Feed forward neural network
- Back propagation algorithm to train neural network

The Perceptron

- A *special* type of artificial neuron
 - Real-valued inputs
 - Binary output
 - Threshold activation function



The Perceptron

- Bias or Threshold?

- They are essentially the same: bias = – threshold

$$y = \begin{cases} 1, & \text{if } \sum_{i=1}^m w_i x_i + b > 0 \\ 0, & \text{otherwise} \end{cases} \quad y = \begin{cases} 1, & \text{if } \sum_{i=1}^m w_i x_i - T > 0 \\ 0, & \text{otherwise} \end{cases}$$

- Simplify notation: let $x_0 = 1$, $b = -T = w_0 = w_0 x_0$

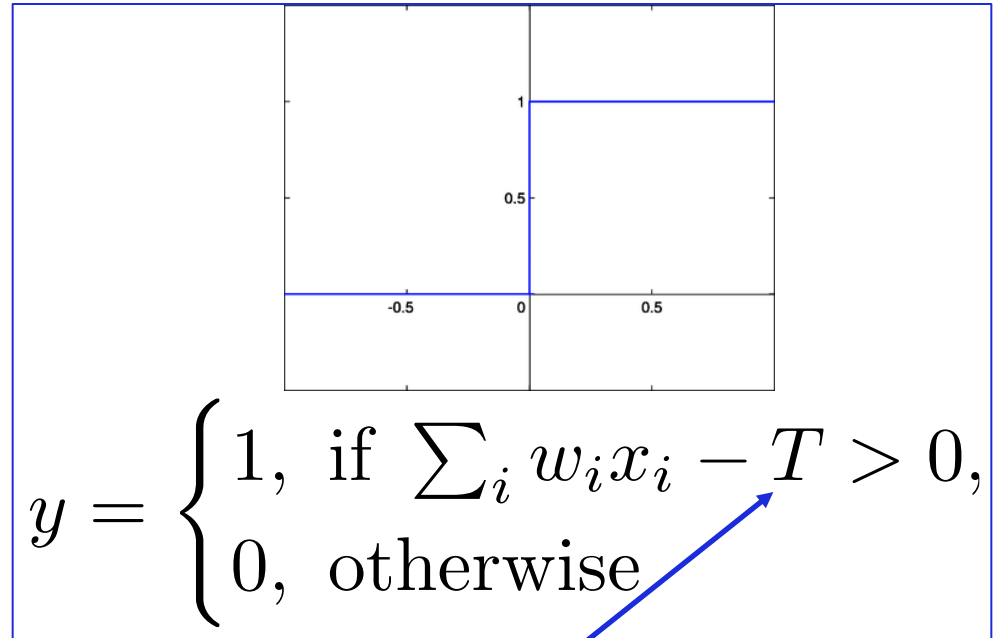

$$y = \begin{cases} 1, & \text{if } \sum_{i=0}^m w_i x_i > 0 \\ 0, & \text{otherwise} \end{cases}$$

- So, we have one block of code for changing all the “weights” rather than changing weights and biases separately

Multi-Layer Perceptron (MLP)


- Add **one hidden node** between the inputs and output

x1	x2	y (class)
0	0	0
1	0	1
0	1	1
1	1	0




Feedforward Neural Network

- A more general form of perceptron
 - Most common type of Artificial Neural Network (ANN)
 - Multiple (hidden) layers, multiple nodes in each layer
 - Each node connects to its adjacent layers
 - Fully connected, NO jump connections
 - A lot of weights: one per link + one bias per node



NNs for (Multi-Class) Classification




How Can We Learn ANN Weights?

- A complex **optimisation** problem!

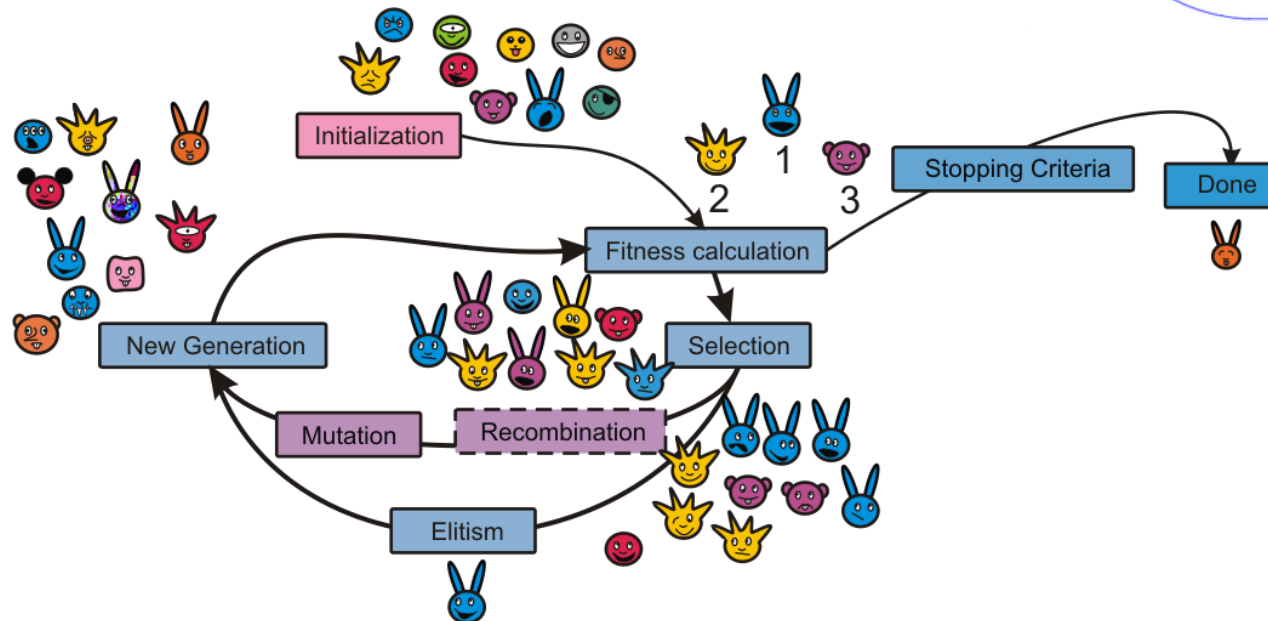

$$\min error = f(w_{ij})$$

- Usually **non-convex** (many local optima)
- Extremely **high-dimensional**
- Not feasible to solve by using exact methods




Learning ANN Weights

- Approximate methods
 - Hill climbing (local search)
 - (Stochastic) gradient descent search
 - Simulated annealing
 - Tabu search
 - Evolutionary computation
 - ...



Training a Neural Network



Training a Neural Network



- **Initialise** the weights (randomly)
- **Feedforward**
 - For each example, calculate the **predicted outputs** o_z using the current weights
 - Calculate the total **error** $\sum_z (d_z - o_z)^2$
- If the error is small enough, we can stop.
- Otherwise, we use **back propagation** to adjust the weights to make the error *smaller*.
 - Uses gradient descent (GD)

Back Propagation (BP) Algorithm

- Estimate the contribution (gradient) of each weight to the *error*, i.e. how much the error will be reduced by changing the weight by the gradient.
- Change each weight (simultaneously) proportional to its contribution to reduce the error as much as possible
 - Move in the direction of the steepest gradient
- We calculate the contribution/gradient backwards (from the last/output layer to the first hidden layer)
- Error of a single **output** node is $d_z - o_z$
 - d_z means “*desired*”
 - o_z means “output” (i.e. what we *actually* got)

Back Propagation (BP) Algorithm


- How **big a change** should we make to **weight $w_{i \rightarrow j}$** ?
 - Make a **big change** if will improve error **a lot** (big contribution)
 - Make a **small change** if **little effect** on error (small contribution)



- β_j is how “**beneficial**” a change is for node j (“error term”)
- When changing $w_{i \rightarrow j}$, the error change should be:
 - Proportional to the **output**: o_i (larger output = more effect)
 - Proportional to the **slope of the activation function** at node j : slope_j
 - Proportional to error term of j (β_j)


Back Propagation (BP) Algorithm

- How to calculate $slope_j$?
 - Some calculus knowledge:
 $derivative$ of the activation function
 - Steeper (larger) the slope, larger the effect of changing the weight
 - We don't expect calculus in this course!



- How to calculate β_j ?
 - $Back-propagated$ from later layer

- The **output layer**: the error $\beta_z = d_z - o_z$
- **Other layers**: error is $\beta_j = \sum_k w_{j \rightarrow k} \times slope_k \times \beta_k$



Back Propagation (BP) Algorithm

- Assume a neural network with:

- Activation function: **sigmoid**

$$slope_j = o_j(1 - o_j)$$

- Target: minimise **total sum squared error**


$$error = \frac{1}{2} \sum_{s \in examples} \sum_{c \in classes} (d_{sc} - o_{sc})^2$$

- Output** node:

$$\beta_z = d_z - o_z$$

- Hidden** node:

$$\beta_j = \sum_k w_{j \rightarrow k} o_k (1 - o_k) \beta_k$$



Makes the
maths easier!
(differentiation)


BP Algorithm Implementation

- Let η be the learning rate (“eta”...)
- Initialise all weights (+bias) to small random values
- Until total error is small enough, repeat:
 - For each input example:
 - Feed forward pass to get predicted outputs
 - Compute $\beta_z = d_z - o_z$ for each output node
 - Compute $\beta_j = \sum_k w_{j \rightarrow k} o_k (1 - o_k) \beta_k$ for each hidden node (working backwards from last to first layer)
 - Compute (+store) the weight changes for all weights
$$\Delta w_{i \rightarrow j} = \eta o_i o_j (1 - o_j) \beta_j$$
(proportional to all 3 factors)
 - Sum up weight changes for all input examples
 - Change weights!

BP Algorithm Example


- Calculate one pass of the BP algorithm given the example (feedforward + back propagation)

Inputs		Outputs	
I_1	I_2	d_5	d_6



Notes on BP Algorithm

- *1 Epoch*: all input examples (entire training set, batch, ...)
- A target of 0 or 1 cannot reasonably be reached. Usually interpret an output > 0.9 or > 0.8 as '1'
- Training may require *thousands* of epochs. A convergence curve will help to decide when to stop (over-fitting?)



Summary

- (Multi-layer) Perceptron
 - Bias and threshold are essential the same
 - Simplify the notation
- Feedforward neural network
- Back propagation
 - Gradient descent
 - Feedforward + error back propagation -> weight change